

Introduction to R package

Thanks to Professor Yan Yu

Why R?

- free !!!
- available for Unix, Linux and Windows
- Nice plots

- Other Statistical Softwares:

SAS (expensive)

S-plus (close relative of R)

Minitab

SPSS

What is R?

- A programming language for statistical data analysis called S developed at Bell Labs. Later extended to S+.
- R is a free, open source version of S/S+.
- Under active cooperative development; versions change frequently.
- Very popular tool in statistics community: New methods often now implemented there.

Installing R

- CRAN: <http://cran.r-project.org/>
 - (R homepage: <http://www.r-project.org/>)
 - Windows 95 and later → Base
 - the latest release (R-2.15.2, 2012-10-26):
<http://cran.case.edu/>
- Statlib <http://lib.stat.cmu.edu/>
- UCI Machine Learning: <http://archive.ics.uci.edu/ml/>
- Nice data mining course notes from Penn State:
<http://sites.stat.psu.edu/~jiali/course/stat557/>

Basics of R

- Command-driven language
- Data (and functions) stored as named objects
- Objects can be fairly simple (vectors, matrices) or more complex (assembled from other objects)
- RStudio(<http://www.rstudio.com/>): Powerful and Free IDE for R
 - Syntax highlighting, code completion, and smart indentation
 - Workspace browser and data viewer
 - Plot history, zooming, and flexible image and PDF export
 - Integrated R help and documentation

Install R, R Studio, R Rattle (clips to follow – thanks to Mai, Feng)

- Install R:
<http://www.youtube.com/watch?v=SJ9sVyqWJn8&hd=1>
- Install RStudio:
<http://www.youtube.com/watch?v=6aTRbo7kdGk&hd=1>
- Install Rattle and Explore Iris:
<http://www.youtube.com/watch?v=7P16fj0tqa4&hd=1>
- Install R and introduction to R (DSI)
<http://www.youtube.com/watch?v=ZmtkqaRVTDc>

Help and documentation

- Use "An introduction to R"
(<http://cran.r-project.org/doc/manuals/R-intro.pdf>)
- Some recommended sections: 1.1, 1.7-1.11, 2, 5.1-5.4, 5.8, 6, 7.1, 10.1, 12.1, 12.3
- `help(mean)`
- `help.search("<subject>")`
- `help.start()`

Organizing your computations

- R has a "current directory" (Working Directory)
 - To set Working Directory in RStudio: Session -> Set Working Directory
- Your objects (loaded datasets, variables, functions, etc.) are contained in your "current workspace", which can be saved any time
 - In Rstudio: Session -> Load Workspace/Save Workspace As
- Keep it tidy!
 - Simple way: Keep separate projects (code, data files) in separate workspaces/directories
 - More advanced: You can use Projects in Rstudio which comes with integration with version control tools such as Git₈ or SVN (<http://www.rstudio.com/ide/docs/using/projects>)

More about R

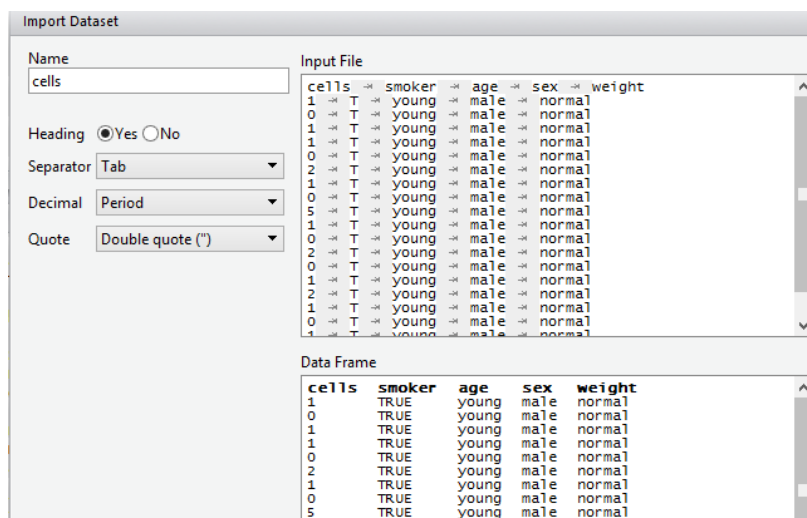

- R workspace
 - Show what's been stored: `ls()` or `objects()`
 - Remove individual objects: `rm(a, b, etc.)`
 - Clear entire workspace: `rm(list=ls())`
 - In RStudio: Session -> Clear Workspace
- Quit: `q()` or file → quit R
- Up arrow repeats prior commands
- R is case sensitive

Assignments and some basic math

- The assignment operator consists of ' $<-$ ' and ' $=$ ' and points from the expression to the name given to that expression
 - $x <- 10$ assigns the number 10 to the letter x
 - $x = 10$ does the same thing
- R as an over-qualified calculator
- \log , \exp
- mean , median , mode , max , min , sd
- Trigonometry: $\cos(x)$
- Set operations: $\text{union}(x, y)$
- Logical operators: $<$, $<=$, $>$, $>=$, $==$, $!=$

Import and export of data

- Useful functions: `read.table`, `write.table`, `scan`
- Works with mixed-type data (numbers and text columns)
- Works well with tab-separated data (connection to Excel)
- Easiest way: Import Dataset in RStudio



Functions

- Most of R consists of functions
- The arguments can either be input in the right order, or using argument names
- Use `help(...)` !
- In Rstudio, hit tab after function name gives help about arguments

Vectors, matrices and data frames

- Calculations often done on vectors, matrices or data frames
 - Vectors: `x = c(10,20,30,40)`
 - Matrix: `A = matrix(c(1,2,3,4), nrow = 2)`
 - Data frames are the usual datasets, with variables names, can have both categorical and numerical variables
 - `data.frame(A)` converts `A` into a dataframe
- Elementwise operations
- Subsetting, indexing
- Logical indexing

Vectors

- Vectors have length, but no dimension
 - `>length(vector)`
- `>c()` concatenates or combines numbers inside `()` into a vector or list
- `>seq()` generates sequence
 - Can specify length of sequence and increment
- `>sort(vector)` sorts items in vector
- Can use usual arithmetic and mathematical functions on vectors
- Logical vectors
- Character vectors
- Indexing a vector:
 - `>vector[i]`
 - `>vector[-i]`

Matrices

- Matrices have length and dimensions
 - `>length(M); dim(M)`
- Generating matrices
 - By combining vectors: `rbind()`; `cbind()`
 - `>matrix()` command
- Transpose
 - `>t(M)`
- Diagonal
 - `>diag(M)`
- Inverse
 - `>solve(M)`
- Multiplying matrices: `M%*%N`
- Indexing matrices

Data Frame

- `>summary(data)`
- `>names(data); >attributes(data)`
- Editing data
 - `>fix(data)` or `>edit(data)`
- `>data$var` select *var* from *data*
 - In RStudio, hit tab after `data$` allows you to select/autocomplete variable names in *data*
 - Need to select more than 1 variable?
 - `>myvars = c('var1', 'var2')`
 - `>data[,myvars]`
 - You can also use numerical index as in matrix
 - `>data[,c(1,3)]`

Lists

- "Compound objects" can often be constructed as lists
- Using \$ to access parts of lists
- data.frame
- names

R packages

- A package: collection of functions (and data) concerning special application
- Contributed from different sources/persons
- Can be downloaded from CRAN or BioConductor
- must be "loaded" with `library()`
- `search()`
- `help(package = graphics)`
- Also: Documentation from `help.start()`

Probability distributions

- norm, binom, beta, cauchy, chisq, exp, f, gamma, geom, hyper, lnorm, logis, nbinom, t, unif, weibull, wilcox
- Four prefixes:
 - 'd' for density (PDF)
`>dbinom(x=4,size=10,prob=0.5)`
 - 'p' for distribution (CDF)
`>pnorm(1.86)`
 - 'q' for quantile (percentiles)
`>qnorm(0.975)`
 - 'r' for random generation (simulation)
`>rnorm(10)`
- Each distribution has arguments that need to be specified
 - `>rnorm(n=10,mean=100,sd=20)`

Load External Files

- Stata, SPSS, SAS files
 - Library(foreign)
 - Stata: read.dta
 - SPSS: read.spss
 - SAS: read.xport (must first create export file in SAS)
- Excel files
 - Files must be saved as comma separated value or .csv
 - read.table, read.csv, read.csv2: identical except for defaults
- Watch the direction of '/'!
- >load(".Rdata")
- Loading and running R programs
 - >source(".R")

NA, NaN, and Null

- NA or “Not Available”
 - Applies to many modes – character, numeric, etc.
 - Missing values are NAs
- NaN or “Not a Number”
 - Applies only to numeric modes
- NULL
 - Lists with zero length

Lm, glm, nlm, optim

■ Linear models

- `>lm(y~x1+x2, data = dataset)`
- Example:
 `>data(cars)`
 `>lm(dist~speed, data=cars)`

■ Generalized linear models

- `>glm(y~x1+x2, family="", data = dataset)`
- See `help(family)` for family options; family specifies error distribution

■ Non-linear minimization

- `>nlm(f, p)` where `f` is the function to be minimized and `p` are the starting parameters

■ Optimization


- `>optim(par, fn)` where `par` are the initial values and `fn` is the function to be minimized

Graphical visualization

- A "generic" function: `plot()`
 - Producing Simple Graphs with R(<http://www.harding.edu/fmccown/r/>)
- High level commands, like `pairs`, `image`, `contour`...
- Lower level commands, adding stuff: `points()`, `lines()`, `text()`, `title()`, `legend()`...
- plotting characters `pch`, colors `col`...
- More advanced packages:
 - `ggplot2` (<http://blog.echen.me/2012/01/17/quick-introduction-to-ggplot2/>)
 - `Lattice` (http://www.isid.ac.in/~deepayan/R-tutorials/labs/04_lattice_lab.pdf)

Graphics Options

- `>par()` is used to set graphical options
 - `mfrow` and `mfcol`
- `>plot(x,y)`
 - `>identify()`; `>abline()`
 - Arguments: type, point style, labels, etc.
- Boxplot, histograms, sunflowerplots, piecharts, etc.
- Saving them
 - `>dev.off()` when things go wrong
 - Know the current directory
 - In RStudio



Writing your own functions

- Collecting commands into a function
- Arguments to function
- Returning result as a list
- Assignments within functions
- programming: conditional statements, loops (be careful, they are slow in R) etc...
- `fix(<myfunction>)` to edit a function

Write simple functions

■ Writing a user function

- `>function(argument list){expressions}`
- Expressions or commands can be grouped together in `{}`, as they often are when writing functions.

- Example:

```
abs_val = function(x){  
  if(x >= 0){  
    return(x)  
  }  
  else{  
    return(-x)  
  }  
}  
abs_val(-5)
```

■ Conditionals

- 'if', 'else'

■ Loops

- 'for'

Help!

- `>help(package=stats)`
 - Lists functions in the specified package
- `>help(glm)` or `>?glm`
- `>help("{")` for characters
- Help → Html help → Search Engine & Keywords
 - Potentially most helpful
- R documentation
 - `command{package}`
- More help:
 - Stackoverflow: <http://stackoverflow.com/questions/tagged/r>
 - R Help Mailing List: <http://r.789695.n4.nabble.com/R-help-f789696.html>

Other Resources

- The R Manuals under Documentation at <http://www.r-project.org>
- Mark Handcock's website at <http://www.stat.washington.edu/handcock/567/links.html>
- Jonathan Barron's R resources at <http://finzi.psych.upenn.edu/>
- Fox, John. 2002. *An R and S-PLUS Companion to Applied Regression*